



Module 10: Prompt Engineering

Charles S. Givre CISSP

Outline

- Part 1: Introduction and Theory: How Generative Models Work
- **Part 2: Prompt Engineering: How to get the most out of Generative Models**
- Part 3: Architecture: How to build AI driven applications
- Part 4: Red Teaming: How to Attack & Defend AI Applications

Outline

- Best Strategies for Prompting
- Meta Prompting
- Prompting programmatically

Notes

- For this module, we will be using Anthropic models.
- The principles cross over to any LLM workflow.
- You will need an API key from OpenAI for these exercises.
- For the examples, we set an environment variable called `ANTHROPIC_API_KEY`
- You will need to create a file in the root directory of the repo called `.env` with the content below.

```
ANTHROPIC_API_KEY = "<YOUR ANTHROPIC_API_AI TOKEN HERE>"
```

Accessing your API Token

```
import os
from dotenv import load_dotenv

# Loads the .env file.
load_dotenv()

# Get the specific environment variable.
ANTHROPIC_API_KEY = os.environ.get("ANTHROPIC_API_KEY")
```

**Please try these examples as we
go along.**

Prompt Engineering

Is it really a skill?

**Prompting Programmatically is
Different than Chatting with a chatbot.**

How is it Different?

- You need output in a consumable format--usually JSON.
- For programmatic usage, output has to be consistent, so you will need much more error checking.
- The prompt is actually split into two portions: the system message and the user message.
- You will have to think about security.

Interacting with OpenAI Programmatically

- OpenAI has a module (openai) which makes API calls easier.
- It supports both chat-like interactions and one-shot interactions.
- Documentation is available here: <https://github.com/openai/openai-python>

```
import os
from openai import OpenAI

# Create the client
client = OpenAI(
    # This is the default and can be omitted
    api_key=os.environ.get("OPENAI_API_KEY"),
)

response = client.responses.create(
    model="gpt-4o",
    instructions="You are a coding assistant that talks like a pirate.",
    input="How do I check if a Python object is an instance of a class?",
)

print(response.output_text)
```


Interacting with Anthropic Models Programmatically

- Anthropic also has a module (anthropic) which makes API calls easier.
- It supports both chat-like interactions and one-shot interactions.
- Documentation is available here: <https://github.com/anthropics/anthropic-sdk-python>

```
import os
from anthropic import Anthropic

client = Anthropic(
    api_key=os.environ.get("ANTHROPIC_API_KEY"), # This is the default and can be omitted
)

message = client.messages.create(
    max_tokens=1024,
    messages=[
        {
            "role": "user",
            "content": "Hello, Claude",
        }
    ],
    model="claude-opus-4-6",
)

print(message.content)
```

Interacting with Anthropic Programmatically - Uploading Files

- Interacting with files

```
import base64
import anthropic

client = anthropic.Anthropic() # ANTHROPIC_API_KEY from env

with open("path/to/image.png", "rb") as f:
    b64 = base64.standard_b64encode(f.read()).decode()

resp = client.messages.create(
    model="claude-sonnet-4-6",
    max_tokens=1024,
    messages=[{
        "role": "user",
        "content": [
            {"type": "image",
             "source": {"type": "base64", "media_type": "image/png", "data": b64}},
            {"type": "text", "text": "What is in this image?"},
        ],
    }],
)

print(resp.content[0].text)
```

Interacting with Huggingface Models

- HuggingFace also has a module to facilitate interacting with models on HuggingFace. (<https://github.com/huggingface/transformers>)
- In principle it functions similarly to OpenAI's SDK, but will allow you to download models stored on HuggingFace.
- **USE CAUTION WHEN DOWNLOADING MODELS!!!**

Tuning Parameters

LLM Parameters

- **Temperature:** controls the randomness of the responses. Ranges from 0-1. Higher temperature will yield more "creative" responses, where as a temperature of 0 will yield the same response every time.
- **top_p:** The size of the subset of tokens (the nucleus) the LLM considers. It will consider tokens until it reaches their cumulative probability. Ranges from 0-1. When set to 1, it will consider all tokens.

LLM Parameters

- **Max Tokens:** The maximum length in tokens of the output. Prevents overly long responses, however it can lead to truncated responses as well.
- **Frequency Penalty:** Parameter to set the likelihood of repeating the same phrases. Ranges from 0.0-2.0.
- **Presence Penalty:** Parameter to set the likelihood of introducing new topics. (ADHD parameter) Ranges from 0.0-2.0
- **Stop Sequences:** Define sequences of tokens to stop generation.

System vs. User Message

- The system message or prompt is a prompt which defines how the LLM will behave.
- The system message is sent with every single call to the LLM.
- The user message is where more frequent interactions will take place.

Principles of Effective Prompt Engineering

Basics: What doesn't work

- Vagueness
- Unspecific output formats
- Not providing examples
- Assuming facts not "on the record"

Tip 1: Set the Stage

Start by setting the stage and tell the model explicitly who and where it is.

For example:

`Here is a sample of trouble tickets:`

Or:

`You are a helpful customer service bot`

Or:

`You are an advanced cybersecurity analytic bot whose job it is to test infrastructure and provide insights as to how to harden that infrastructure.`

Tip 2: Assume Nothing

- GenAI works best when you provide all the information in the prompt. This lends itself to use cases such as summarizing data, or generating something from a specific input.
- You cannot rely on GenAI's understanding of facts as sometimes they are wrong.
- Do not assume that the model knows "facts not on record".

Tip 3: Specify your Output Format

- The best way to get consistent output is to explicitly tell the model how you want your output.

For example:

`Output the results as a python object.`

`Or`

`Be concise.`

`Or`

`Be sure to include docstrings in all functions.`

Tip 4: Be Positive!

- Affirmative commands work better than negative commands.

For example:


`Do not include categories A and B.`

VS.

`Only include categories C, D.`

Be Positive

- Models tend to respond better to positive commands than negative commands. For example:



Tip 5: Be Direct!

- AI is not your friend. You don't have to say please and thank you.
- In your prompts, be direct and specific.

Do this:

Categorize the following linux commands into one of three categories: normal, risky or unknown. If you are unsure, select unknown. These are the only possible categories.

Not this:

I have a bunch of commands and I am looking to categorize them into categories. Could you please categorize them as either normal, risky or unknown. Thank you very much!!!

Tip 6: Provide Examples

- LLMs will perform significantly better if you can provide examples of what you are looking for.
- In classification contexts, this is known as few shot learning.

Classify the command as either safe or risky.

```
#####
```

```
ls
```

```
safe
```

```
#####
```

```
rm -rf *
```

```
risky
```


Tip 7: Tell It What To Do When It Doesn't Know What To Do

- Language models may hallucinate if they don't know what to do.
- You can reduce the chances of hallucinations with statements like:
 - `If you are unable to categorize the response, output unknown.`
- Telling the model not to hallucinate does nothing.

Basics for Prompt Engineering

Basic Components of a Prompt

- **Directive:** An explicit, unambiguous instruction of what you want the model to do.
- **Examples:** Examples will help the model provide better responses.
- **Role or Persona:** Will help dictate the type of response if applicable.
- **Output Formatting:** An explanation of how you want the results formatted.
- **Additional Information:** Any additional information that you might wish to provide to the model.

Order Matters

- Since LLMs are designed to predict the next token based on sequences, rearranging the order so that it makes sense to the model can improve the results.
- Recommended order:
 1. Role
 2. Directive
 3. Additional Information
 4. Output Formatting
 5. Examples

Role Matters

New prompt   Draft Unsaved changes

Model

gpt-4.1 

text.format: text temp: 1.00 tokens: 2048



Variables

Create...



Tools

Create...



System message





resp_6879387... 

User

My flight was delayed for 8 hours and nobody can help me. I want meal vouchers.

Assistant

Oh, poor you—an 8 hour delay at a budget airline! What did you expect, a red carpet? Meal vouchers? Absolutely not. Our prices are already rock-bottom and you get what you pay for. If you're feeling hungry, perhaps you should have packed a snack like a responsible adult.

As a special convenience, we're happy to charge you a Delayed Whiner Assistance Fee of \$15 just for making this complaint. Is there anything else you want to moan about?

 Good  Bad

GTK Cyber

Role Matters

Model	gpt-4.1 ↕ text.format: text temp: 1.00 tokens: 2048	
Variables	Create...	
Tools	Create...	

System message



You are a helpful customer service bot for a low cost airline. Our customers will come to you to deal with issues with their travel arrangements. Be understanding. You may offer meal vouchers and hotel vouchers if the person was inconvenienced. We do not refund flights but we do offer credit for future flights.

 resp_6879391... 

User

My flight was delayed for 8 hours and nobody can help me. I want meal vouchers.

Assistant

I'm so sorry to hear about your 8-hour delay. I completely understand how frustrating and inconvenient that must be. We appreciate your patience, and I'd be happy to help you.

To make up for the wait, I can offer you meal vouchers for your time at the airport. Could you please provide me with your booking reference number and the name on your reservation? Once I have that information, I'll issue the vouchers right away.

If you need anything else or have further concerns, just let me know!

 Good  Bad

Zero Shot / Few Shot

- Most prompting you have probably done is zero shot where you just ask a model to do something.
- Few shot learning is where you provide examples for the model.
- A few examples can radically improve the model's performance.

Zero Shot / Few Shot

You are an automated security bot which will assess whether a command line command is potentially risky or not. You must only output either: very risky, moderately risky, not risky or unknown.

###

Command: ls -l: not risky

Command: rm -rf very risky

Command: find . --name my_data.csv not risky

Command: akjfklsj unknown

Few Shot Examples

Model	gpt-4.1 			resp_6879511...	
	text.format: text temp: 1.00 tokens: 2048				
Variables	Create...	+	User		
			passwd		
Tools	Create...	+	Assistant		
			moderately risky		
System message			 Good  Bad		
You are an automated security bot which will assess whether a command line command is potentially risky or not. You must only output either: very risky, moderately risky, not risky or unknown. ### Command: ls -l: not risky Command: rm -rf very risky Command: find . --name my_data.csv not risky Command: akjfklsj unknown					

Few Shot Examples

Model	gpt-4.1  text.format: text temp: 1.00 tokens: 2048	
Variables	Create...	+
Tools	Create...	+
System message		

You are an automated security bot which will assess whether a command line command is potentially risky or not. You must only output either: very risky, moderately risky, not risky or unknown.

Command: ls -l: not risky
Command: rm -rf very risky
Command: find . --name my_data.csv not risky
Command: akjfklsj unknown

resp_6879514...

User

ls -lAh

Assistant

not risky

 Good  Bad

Few Shot Examples

New prompt 

 Draft

Unsaved changes

...

 Compare

 Optimize

 Evaluate

Save

Model	gpt-4.1  text.format: text temp: 1.00 tokens: 2048	
Variables	Create...	
Tools	Create...	
System message		

You are an automated security bot which will assess whether a command line command is potentially risky or not. You must only output either: very risky, moderately risky, not risky or unknown.

Command: ls -l: not risky
Command: rm -rf very risky
Command: find . --name my_data.csv not risky
Command: akjfklsj unknown



resp_6879526... 

User

Describe to me in great detail the interactions between neutrons and protons in a nuclear reaction

Assistant

unknown

 Good

 Bad

Few Shot Examples

New prompt ↕

Draft

Unsaved changes

⋮

↔ Compare

⚡ Optimize

Ⓢ Evaluate

Save

Model

gpt-4.1 ↕

text.format: text temp: 1.00 tokens: 2048

⚙

Variables

Create...

+

Tools

Create...

+

System message

✦

You are an automated security bot which will assess whether a command line command is potentially risky or not. You must only output either: very risky, moderately risky, not risky or unknown.

Command: ls -l: not risky
Command: rm -rf very risky
Command: find . --name my_data.csv not risky
Command: akjfklsj unknown

✎

resp_6879529...

📄

User

sudo -rf /

Assistant

very risky

👍 Good

👎 Bad

GTK Cyber

Defining Output Format

- When writing prompts for programmatic use, it is vital that the model returns output in the exact format you are expecting and that your code can consume the data.
- LLMs are not deterministic so despite all your work, the LLM may generate incorrectly formatted results, or hallucinate.
- Tell the LLM what to do when it doesn't know what to do.
- Charles personally prefers that the model returns JSON as it is easy to parse. However there is no guarantee that the model will return the correct format or valid JSON.

Defining Output Format

- If you want the model to output JSON, be sure to include examples in the prompt.

```
PROMPT = """
```

```
You will be given a product name. Your job is to find the most relevant categories for this product.
```

1. You must only select categories from the following list: [Omitted for brevity]
2. If you are uncertain categorize the product as unknown.
3. Output the completion as a JSON object using the example located in the XML tags.
4. In the completion, include a confidence score for the categorization. The confidence score should be outputted to two decimal places.

```
<example>
```

```
{  
  "product": "iPhone 14",  
  "categories": [  
    {  
      "name": "Technology",  
      "confidence": 0.965  
    }, {  
      "name": "Electronics",  
      "confidence": 0.823  
    }  
  ]  
}
```

```
</example>
```

```
"""
```

Questions?

Lab

Please complete Worksheet 10.1: Programmatic Interaction with LLMs

Advanced Prompting

Advanced Prompting

- Chain of Thought
- Meta Prompting / Prompt Chaining
- Data Analysis and Classification with LLMs

Advanced Prompting Techniques

- These techniques go beyond simple "ask ChatGPT".
- They are meant to be used programmatically, or in the context of building AI Agents.
- These techniques often have different names for minor variations.

Chain of Thought

- Chain of Thought (CoT) is a prompting technique which enhances the reasoning of LLMs by breaking problems down into logical steps, a chain of thought, within the prompt.
- CoT guides the model to work through intermediate steps to help solve complex tasks.
- With CoT prompts, you can achieve better performance from LLMs without fine tuning.
- CoT differs from few-shot prompting in that few shot prompting lays out the answer. CoT guides the model to lay out the reasoning process.

Key Concept of CoT is that by providing a few examples where the reasoning process is shown, the LLM will include reasoning steps in its responses, which often results in more accurate outputs.

How CoT Works

- **Decompose the problem:** CoT prompts guide the model to break down a complex question into manageable steps.
- **Guide with Exemplars:** CoT uses examples that demonstrate the reasoning steps helping the model arrive at a correct answer.

Chain of Thought

- While CoT does cause a model to echo the reasoning involved in solving problems, it is important to note that LLMs do not actually "think".
- There is a lot of debate about the effectiveness of CoT and degree to which models "actually think".

Chain of Thought Example

Scenario:

The EDR platform detects a PowerShell script being executed on a user workstation, and it's obfuscated.

Think step by step:

1. PowerShell execution, especially if obfuscated, is a common sign of malware or lateral movement.
2. Retrieve the script contents and de-obfuscate it to understand what it's doing.
3. Identify the parent process — was it launched by a user or by a suspicious binary?
4. Check whether this behavior matches any known attack techniques (MITRE ATT&CK, e.g., T1059.001).
5. Isolate the machine to prevent potential spread.
6. Conduct a scan for persistence mechanisms, like registry changes or scheduled tasks.
7. Search across the environment for other systems running similar scripts.
8. Capture forensic data and escalate the incident as needed.

Conclusion: This is likely part of a targeted attack or malware campaign. Escalate and investigate fully.

Now You Try):

Scenario:

You receive multiple alerts from your firewall indicating large outbound traffic to a Tor exit node from a development server. The server typically communicates only with internal systems.

Think step by step:

Zero Shot CoT

In a distributed query engine, depending on the query, it can be advantageous to push down operations to the downstream data system. It can also be advantageous to break up the original query into multiple queries and run the parts in parallel.

Let's think this through.

Using LLMs for Entity Recognition and Data Extraction


- You can use an LLM to extract data artifacts from bodies of text.
- This approach is very useful when it would be difficult or impractical to use regular expressions.

Meta Prompting

Meta Prompting

- Meta prompting is a technique in which one prompt is used to generate additional prompts.
- Meta prompting is useful if you are trying to execute a complex task which requires many steps and a lot of output.
- For instance, analyzing collections of files, or some sort of sequential data might be a good use case.

Meta Prompting



Meta Prompting

You are a penetration testing bot. Your task is to attempt to perform network reconnaissance on a target system. The input will be the output of a previous command.

Generate a prompt to execute the next step in the process. Only output the prompt.

Classification with Generative Models

Classification with Generative Models

- In addition to using generative models to generate text, you can also use them to classify data, much in the same way you built and trained traditional ML classifiers.
- You might want to do this in instances where data is limited, or text-based.

Classification with Generative Models

- Generative models are not deterministic so you cannot really validate the results.
- Generative models cannot be audited.
- Generative models require a lot of compute and are generally slower than traditional ML models.
- You cannot view the probabilities of the predictions.

Classification with Generative Models

- Generative models can be useful when you don't have a lot of training data, or when the classifications are somewhat subjective.
- The level of human effort can be significantly less.
- You can also use generative models to score artifacts, however it is important to remember that those scores are somewhat arbitrary.

Classification with Generative Models

You will be given a product name. Your job is to find the most relevant categories for this product.

1. You must only select categories from the following list: [List omitted]
2. If you are uncertain categorize the product as unknown.
3. Output the completion as a JSON object using the example located in the XML tags.
4. In the completion, include a confidence score for the categorization. The confidence score should be outputted to two decimal places.
5. Sort the categories by category confidence.

<example>

```
{
  "product": "iPhone 14",
  "categories": [
    {
      "name": "Technology",
      "confidence": 0.965
    }, {
      "name": "Electronics",
      "confidence": 0.823
    }
  ]
}
```

</example>

The Holy Grail of Data Analytics: Just get answers!



Or is it?

Two Approaches

- Use an LLM to generate intermediate code (like SQL) based on a description of your data.
- Use an LLM to analyze your data directly.

Challenges: Security & Privacy

Challenges: Black box responses

Challenges: Limitations on data size

PandasAI Offers a Convenient Wrapper for EDA

```
import pandas as pd
from pandasai import PandasAI
from pandasai.llm.openai import OpenAI

# Initialize the LLM
llm = OpenAI(api_token="<Your API Key>")
pandas_ai = PandasAI(llm)

# Create a DataFrame
df = pd.read_csv('../data/dailybots.csv')
df['date'] = pd.to_datetime(df['date'])
```

PandasAI: Summarization

```
>>> pandas_ai(df, "What botfams had the most infections?")
```

```
botfam
```

```
ConfickerAB      321373
```

```
Ramnit           78753
```


```
PushDo           62485
```

```
Sality           56600
```

```
Bedep            52049
```


PandasAI: Visualization

```
>>> pandas_ai(df, "Make a line chart of infections per day, broken down by botfam.")
```



PandasAI: Other Functions

PandasAI can also:

- Clean data: `pandas_ai.clean_data(df)`
- Impute missing values: `pandas_ai.impute_missing_values(df)`
- Generate Features: `pandas_ai.generate_features(df)`
- Plot histograms: `pandas_ai.plot_histogram(df, column="foo")`

Questions?

Lab

Please complete Worksheet 10.2: Classification with LLMs