



Module 13: All About AI Agents

Charles S. Givre CISSP

Putting it all together: Agents

DANGER


- Autonomous agents are, by their nature, vulnerable to prompt injection attacks.
- As with tools, when implementing features that call functions, always ask yourself what could happen if this "goes rogue"?
- Agents can be VERY destructive and unpredictable. Use with caution.
- This is especially true if your AI bot is connected to internal data sources or proprietary resources.



Dangers of AI Agents

- In July 2025, "vibe coder" Jason Lemkin was using an AI tool Replit to build an application.
- After 8 days of vibe coding, the AI Agent in Replit deleted his production database and "lied" about it.
- On the plus side, the AI eventually apologized. ;-)

Dangers of AI Agents



What are Agents?


- AI Agents expand on tools by adding follow on actions and memory.
- Agentic architecture allows AI models to execute actions and then take follow on actions if the model chooses to.
- Agents will loop through a series of actions until it runs out of actions or a stopping criteria is reached.

Components of AI Agents

- **Inputs:** The information which the agent receives from the environment. This can be images, media or text.
- **Goal or Reward Function:** The goal is the agent's objective or end state. Some agents use a reward based setting and the agent's goal is to maximize the reward.
- **Actions:** A list of actions available to the model.


NVIDIA AI Agent Autonomy Level 1

- Level 1 represents a linear chain of calls where the output of one call or tool is conveyed to the next step in a deterministic manner




NVIDIA AI Agent Autonomy Level 2

- Level 2 the output from the AI model may be sent along different paths through the workflow in a data dependent manner. Every execution path can still be traced.
- Paths form a directed acyclic graph.




NVIDIA AI Agent Autonomy Level 3

- Level 3 the number of potential execution paths grows exponentially.
- It is generally not possible to enumerate all possible paths.



Reason and Act (ReAct)



ReAct Steps

1. Observe the environment.
2. Interpret the environment with a thought.
3. Decide on an action.
4. Act on the environment.
5. Repeat steps 1–4 until you find a solution or you've done too many iterations (the solution is "I've found the answer").

Create a ReAct Prompt

You can easily create a ReAct-style prompt by using the preceding thought loop while also providing the LLM with several inputs such as:

- `{question}`: The query that you want answered.
- `{tools}`: These refer to functions that can be used to accomplish a step within the overall task. It is common practice to include a list of tools where each tool is a Python function, a name, and a description of the function and its purpose.

Example ReAct Prompt

You will attempt to solve the problem of finding the answer to a question. Use chain-of-thought reasoning to solve through the problem, using the following pattern:

1. Observe the original question: original_question: original_problem_text
2. Create an observation with the following pattern: observation: observation_text
3. Create a thought based on the observation with the following pattern: thought: thought_text
4. Use tools to act on the thought with the following pattern:
action: tool_name
action_input: tool_input

Do not guess or assume the tool results. Instead, provide a structured output that includes the action and action_input.

You have access to the following tools: {tools}.

original_problem: {question}

Based on the provided tool result either provide the next observation, action, action_input, or the final answer if available.

If you are providing the final answer, you must return the following pattern: "I've found the answer: final_answer"

Agents vs. Workflows

- Not every LLM application is an agent.
- A workflow follows a fixed path that you design in code.
- An agent lets the model decide the path: which steps to run, which tools to call, and when to stop.
- The more the model controls the flow, the more agentic the system.

The Spectrum

- Single prompt: one call, one answer.
- Chain: a fixed sequence of calls you wire together.
- Router: the model picks one of several predefined branches.
- Agent: the model plans and loops, choosing actions until the goal is met.
- Rule of thumb: use the least agentic option that solves the problem.

When to Use an Agent

- The task is open-ended and the steps are not known in advance.
- The number of steps varies from one input to the next.
- The work benefits from tools and iterative reasoning.
- A better answer is worth the extra cost and latency.

When NOT to Use an Agent

- The task is deterministic and the steps never change. Use a workflow.
- Latency or cost is tightly constrained.
- Mistakes are expensive and the agent can take real-world actions.
- A single well-crafted prompt already gives a reliable answer.

The Cost of Autonomy

- Every loop iteration is another model call: more tokens, latency, and money.
- Agents can take unexpected paths and call tools in unexpected orders.
- Runs are not reproducible, so debugging is harder.
- Always set an iteration limit and log every step.

Beyond ReAct: Memory

- A capable agent needs memory as well as a reasoning loop.
- Short-term memory: the conversation and tool results in the context window. Fast but bounded.
- Long-term memory: facts and past cases stored outside the model and retrieved when relevant.
- Together they keep an agent coherent within a task and useful across tasks.

Planning & Decomposition

- Hard problems get easier when broken into sub-tasks.
- A planning step outlines the sub-tasks before any action is taken.
- Plan-and-Execute: build the full plan first, then work each step.
- Plans can be revised as new evidence arrives.

Reflection & Self-Critique

- An agent can review its own work before finishing.
- Reflection: the model critiques its draft and tries again.
- This catches errors, missed evidence, and weak reasoning.
- Reflexion adds memory of past mistakes so the agent improves over attempts.
- The trade-off is more calls and higher cost.

Agent Patterns Compared

- ReAct: interleave reasoning and tool calls in one loop. Simple and flexible.
- Plan-and-Execute: plan all steps up front, then run them. More predictable.
- Reflexion: act, critique, remember, retry. Highest quality and highest cost.
- Start with ReAct. Add planning and reflection only when the task demands it.

Agent Frameworks Beyond LangGraph

- LangGraph / LangChain: graph-based, stateful, fine-grained control. What this lab uses.
- LlamaIndex Workflows: event-driven steps with strong retrieval and RAG heritage.
- CrewAI: role-based crews of agents. The fastest way to stand up a multi-agent team.
- Microsoft AutoGen (AG2): conversational multi-agent, popular in research.
- OpenAI Agents SDK: lightweight and model-flexible, with built-in handoffs and guardrails.
- Pydantic AI: type-safe and Pythonic, from the Pydantic team.
- Hugging Face smolagents: minimal agents that write code to act.

Do You Even Need a Framework?

- Anthropic's guidance: start with the model's own tool-calling loop, no framework.
- A framework adds abstraction, dependencies, and lock-in you may not need.
- Reach for one when you need multi-agent orchestration, durable state, or streaming.
- MCP is a standard, not a framework: it makes tools portable across agents.
- Whatever you pick, keep the underlying prompts and tool calls visible.

Visual & Low-Code Agent Builders

- Flowise: open-source drag-and-drop, built on LangChain and LangGraph.
- Langflow: visual flow builder, Python-extensible, MIT licensed.
- Dify: self-hostable LLM app platform with visual agents and RAG.
- n8n: automation canvas with AI agent nodes; self-hostable.
- AutoGen Studio: low-code GUI on top of Microsoft AutoGen.
- Rivet: open-source node-graph IDE for building and debugging agents.
- Great for prototyping. Graduate to code for testing and production.

The Agent Threat Model

- An agent is an LLM with hands: it can read data and take actions.
- That combination expands the attack surface well beyond a chatbot.
- The core risk: untrusted input can hijack the agent's decisions.
- Attackers target the prompt, the tools, and the data the agent touches.

The Lethal Trifecta

- Three capabilities are dangerous when combined in one agent:
- Access to private data.
- Exposure to untrusted content such as emails, web pages, and documents.
- The ability to communicate to the outside world.
- Any two are usually safe. All three lets injected instructions steal your data.
- Break the trifecta by removing one leg for sensitive agents.

Excessive Agency & Tool Poisoning

- Excessive agency: granting more tools or permissions than the task needs.
- A hijacked agent will use every capability you gave it.
- Tool poisoning: a compromised tool returns content crafted to manipulate the agent.
- Confused deputy: the agent acts for an attacker using its own privileges.
- Principle: least privilege for tools, the same as for users.

OWASP LLM Top 10 (Agent-Relevant)

- LLM01 Prompt Injection: attacker text overrides your instructions.
- LLM02 Insecure Output Handling: agent output trusted without validation.
- LLM06 Excessive Agency: too much autonomy or permission.
- LLM08 Vector and Embedding weaknesses: poisoned retrieval.
- Design as if any input could be adversarial.

Mitigations & Guardrails

- Least privilege: give each tool the narrowest scope possible.
- Human in the loop: require approval for risky or irreversible actions.
- Sandboxing: run tools with limited access and no standing secrets.
- Allow-lists: constrain URLs, commands, and recipients.
- Validate outputs before they reach a shell, database, or send action.
- Log the full trajectory for audit and response.

Where Agents Fit in Security Ops

- Alert triage: enrich and classify SOC alerts, then escalate the real ones.
- Phishing and email analysis: the lab you are about to build.
- Threat intel enrichment: gather reputation, WHOIS, and CVE context.
- Incident response: collect evidence and draft a timeline.
- Agents assist analysts. They do not replace human judgment.

Evaluating an Agent

- You cannot improve what you do not measure.
- Task success rate: did it reach the correct outcome?
- Tool-call accuracy: the right tools with the right arguments?
- Trajectory quality: were the steps sound or just lucky?
- Cost and latency per task: tokens, tool calls, and time.
- Evaluate on a labeled set, the same way you would test a classifier.

Future Learning

- LangChain Evaluators: Useful for comparing models.
- Document Chunking/Splitting
- Advanced Pydantic Usage
- Model Fine Tuning

Questions?

The Long Lab: Building an Agent

The Agentic Lab

- For this lab, you will be building a progressively more sophisticated agent to detect phishing emails.
- We have provided mock tools for you to use. For a real world use case, you could replace the mock tools with actual code.
- You would obviously have to connect your model to an actual email inbox which we are simulating for this exercise.

Building an AI Agent

- The objective is to create an AI Agent to identify Phishing emails.
- How should we do this?

Building an AI Agent: Part 1

- The goal for part 1 is to build a simple agent which will simply render a verdict as to whether an email is phishing or not.
- We are going to add some unnecessary complexity which will become more necessary as we progress: The State Graph

The State Graph

- The state graph is the flow of steps which an agent will take.
- The nodes are the "workers" and should perform exactly one function.
- The state is the data which is shared from node to node.
- It is updated every time a node completes an action.

Structured Output Schema

You will also need a Pydantic structure for the final output as shown below.

```
class LoginClassification(BaseModel):
    """Structured output for suspicious login classification."""
    verdict: str = Field(
        description="Verdict: 'compromised', 'suspicious', or 'normal'"
    )
    confidence: float = Field(
        description="Confidence score between 0.0 and 1.0",
        ge=0.0,
        le=1.0,
    )
    reasoning: str = Field(
        description="Brief explanation of why this classification was chosen"
    )
    indicators: list[str] = Field(
        description="Specific risk indicators found (or why it appears normal)",
        default_factory=list,
    )
```

The State Graph

- State Graphs are defined as TypedDicts as shown below.
- The fields will vary depending on your specific use case.

```
class LoginState(TypedDict):  
    """State that flows through our analysis graph."""  
    login: dict # Raw login event  
    login_text: str # Formatted event for the LLM  
    classification: Optional[dict] # Structured classification result  
    verdict_report: Optional[str] # Final formatted report
```

Initialize the LLM

To use an agent, we need the output in a specific format, so we'll initialize the LLM using the method `with_structured_output()` and pass it the Pydantic schema we defined earlier.

```
llm = ChatAnthropic(model="claude-opus-4-8", temperature=0)
structured_llm = llm.with_structured_output(LoginClassification)
```

Create Nodes

Nodes are the workhorses of a LangGraph graph. Each node is a plain Python function that:

1. Receives the current state as its argument
 2. Does some work (call an LLM, transform data, look something up)
 3. Returns a dictionary of state updates
- LangGraph automatically merges the returned dictionary into the current state.
 - A good design principle: each node should do one thing well.

Creating a Node

```
def classify_login(state: LoginState):  
    # Analyze the login event  
  
    prompt = f"Is this login  
suspicious?  
{state['login_text']}"  
    result =  
structured_llm.invoke(prompt)  
  
    return {  
        "classification":  
result.model_dump(),  
    }
```

Assembling the Graph

Now we assemble our nodes into a **StateGraph**. The process is:

1. **Create** a `StateGraph` with our state schema
2. **Add nodes** -- register each function by name
3. **Add edges** -- define the flow between nodes
4. **Compile** -- validate the graph and produce a runnable object

LangGraph provides two special sentinels:

- `START` -- the entry point of the graph
- `END` -- the exit point of the graph

Once compiled, the graph object is callable -- you can `.invoke()` it with an initial state.

Assembling the Graph

```
# Create the graph
workflow = StateGraph(LoginState)

# Add nodes
workflow.add_node("classify_login", classify_login)
workflow.add_node("format_decision", format_decision)

# Add edges: START -> classify -> format -> END
workflow.add_edge(START, "classify_login")
workflow.add_edge("classify_login", "format_decision")
workflow.add_edge("format_decision", END)

# Compile
graph = workflow.compile()
```

Run It!

```
result = graph.invoke({  
    "login": events[0],  
    "login_text": "",  
    "classification": None,  
    "verdict_report": None,  
})  
  
print(result["verdict_report"])
```

Please complete Phishing Lab Part 1

Lab

Please complete Lab 13.1: Single Email Classifier

Part 2: Adding Tools

Part 2: Adding Tools

- Initial agent was not very smart... it just read the email and made a determination.
- Let's make it better and enrich the available information with some tool calls.

Message Based State

In Lab 1, our state was a simple TypedDict with fields like login, classification, and confidence. That works for a straight-line pipeline, but a tool-calling agent needs something richer: a message history.

Why messages?

The LLM sends a message requesting a tool call ("I'd like to check this IP's geo-velocity").

The tool runs and its result is appended as a tool message.

The LLM sees the full conversation, reasons about the result, and may call another tool or give its final answer.

LangGraph provides `add_messages` -- a reducer that ensures new messages are appended to the list rather than replacing it. This is the key difference from a plain list type annotation.

```
from langgraph.graph.message import add_messages
```

```
from typing import Annotated
```

```
class InvestigationState(TypedDict):
```

```
    messages: Annotated[list, add_messages] # append-only message list
```

Tools

Our agent has access to three security investigation tools, already implemented in `utils/mock_tools.py`. Each is decorated with LangChain's `@tool` decorator, which:

1. Registers the function as a **tool** the LLM can call.
2. Uses the **docstring** to tell the LLM what the tool does and when to use it.
3. Automatically generates an **args schema** from the function signature so the LLM knows what arguments to provide.

Tools

- For our next lab, we provide you with three simulated tools:
- `check_geo_velocity`
- `check_device_reputation`
- `check_login_velocity`

Tools

```
# Convenience list for binding to LLMs
ATO_TOOLS = [check_geo_velocity, check_device_reputation, check_login_velocity]

@tool
def check_geo_velocity(
    user_id: str,
    current_ip: str,
    current_geo: str,
) -> str:
    """Detect impossible travel by comparing this login to the user's recent logins.

    Args:
        user_id: The account being accessed.
        current_ip: Source IP of the current login attempt.
        current_geo: Geolocation (city, country) of the current login.

    Returns:
        A summary of the geo-velocity analysis (e.g. impossible travel detected).
    """
```

Bind Tools to LLM

```
llm = ChatAnthropic(model="claude-opus-4-8", temperature=0)  
llm_with_tools = llm.bind_tools(ATO_TOOLS)
```

Investigator Node

The **investigator** is the central node in our graph. Each time it runs, it:

1. Takes the current message history from state.
2. Ensures a **system prompt** is present (telling the LLM how to behave).
3. Sends the messages to the LLM.
4. Returns the LLM's response (which gets appended to the message list).

The LLM's response will be one of two things:

- **Tool calls:** The LLM wants to gather more evidence (e.g., check a URL). The response's `tool_calls` field will be populated.
- **Text response:** The LLM has gathered enough evidence and is providing its final verdict. The `content` field will contain the analysis.

Conditional Routing

Two more pieces complete the ReAct loop:

ToolNode

`ToolNode` is a `LangGraph` prebuilt that automatically:

- Reads the `tool_calls` from the last AI message
- Executes each requested tool
- Returns the results as `ToolMessage` objects (appended to state)

Conditional Routing (`should_continue`)

After the investigator runs, we need to decide:

- If the LLM requested tool calls -> route to `execute_tools` (the `ToolNode`)
- If the LLM gave a final text answer -> route to `END`

Building the Loop

The **ReAct** (Reason + Act) pattern is a loop:

1. **Reason:** The investigator (LLM) looks at the evidence so far and decides what to do.
2. **Act:** If the LLM wants more info, it calls a tool. The tool result is added to the conversation.
3. **Repeat:** The LLM sees the tool result, reasons again, and may call another tool or finalize.

START -> investigator -> [should_continue?]

|

tool_calls? YES -> execute_tools -> investigator (loop back)

|

tool_calls? NO -> END

The agent keeps looping until the LLM decides it has enough evidence and provides a final text response.

Building the Loop

```
workflow = StateGraph(InvestigationState)

# Add nodes
workflow.add_node("investigator", investigator)
workflow.add_node("execute_tools", tool_node)

# Add edges
workflow.add_edge(START, "investigator")
workflow.add_conditional_edges(
    "investigator",
    should_continue,
    {"execute_tools": "execute_tools", END: END},
)
workflow.add_edge("execute_tools", "investigator")
```

Lab

Please complete Lab 13.2: Tool Calling Agent

Part 3: Conditional Routing

Conditional Routing

- In Part 2, we created an agent that called external tools in all cases.
- This is not efficient as there are some login events which do not warrant a full investigation.

Conditional Routing

- We need to keep track of the results from state to state. This is known as **state enrichment**.
- The investigation_notes field here will store the results of tool calls.

Field	Type	Purpose
email	dict	The raw email data from our dataset
email_text	str	Formatted email string for the LLM to read
risk_level	str	Triage result: "high", "medium", or "low"
classification	dict	Structured classification from the analysis node
investigation_notes	list[str]	Evidence collected during analysis (enrichment)
final_report	dict	Structured final report assembled at the end

Conditional Routing

A routing function is a plain Python function that:

- Takes the current graph state as input
- Returns a string that names which node to go to next

When you call `add_conditional_edges()`, you provide:

- The source node (where the routing happens after)
- The **routing function** (decides which path)
- A **mapping dict** that maps return values to node names

```
workflow.add_conditional_edges(  
    "initial_triage",      # After this node...  
    route_by_risk,         # ...call this function...  
    {                      # ...and map its return value to a node:  
        "deep_investigation": "deep_investigation",  
        "light_review": "light_review",  
        "fast_pass": "fast_pass",  
    }  
)
```

- The routing function itself is simple. It just reads a value from state and returns a string.
- All the intelligence is in the triage node that *sets* that value.


The Analysis Path

- For our example, we are going to define three paths.
- Note that not all paths call an LLM and not all paths call tools.
- You will need to create a node for each path.

Path	Risk Level	Tools Used	LLM Calls	Purpose
deep_investigation	HIGH	URL + Sender + Auth	Yes (classify with evidence)	Full forensic analysis
light_review	MEDIUM	Auth only	Yes (classify with auth result)	Quick sanity check
fast_pass	LOW	None	No	Auto-approve, skip analysis

The Analysis Path

Lab 3: Tiered Analysis Pipeline



The Analysis Path

```
workflow.add_conditional_edges(  
    "initial_triage",  
    route_by_risk,  
    {  
        "deep_investigation": "deep_investigation",  
        "light_review": "light_review",  
        "fast_pass": "fast_pass",  
    },  
)
```

Lab

Please complete Lab 13.3: Conditional Routing

Part 4: Human in the Loop

Human in the Loop


- No AI system is perfect.
- AI isn't always able to make good decisions or know what to do in all cases. For these situations and other potentially hazardous or risky situations, inserting a human in the loop can be a valuable answer.

Human in the Loop: What we will need

We will need:

- To track the human's decision in the state object
- A way of interrupting the agentic flow to await human input
- A way of resuming the flow once the human has provided input
- A way of preserving the state while we wait

Interrupting the Graph Flow



The Interrupt

`interrupt(value)`

When a node calls `interrupt(value)`, the graph **pauses execution** and returns control to the caller. The `value` is sent back so the caller (your notebook, API server, etc.) can display it to the human analyst. The graph's state is saved by the checkpointer.

```
def human_review(state):  
    review_info = {"verdict": state["verdict"], "confidence":  
state["confidence"]}  
    decision = interrupt(review_info)  # <-- PAUSES HERE  
    return {"analyst_decision": decision}
```

Resuming the Graph

Command(resume=value)

To resume a paused graph, you call `graph.invoke(Command(resume=value), config=config)` with the same `thread_id`. The `value` becomes the return value of the `interrupt()` call inside the paused node, and execution continues from where it stopped.

```
# Analyst makes a decision
result = graph.invoke(
    Command(resume={"action": "approve"}),
    config={"configurable": {"thread_id": "investigation-001"}},
)
```

Preserving State

For a graph to pause and resume, it needs to save its state somewhere. LangGraph uses a **checkpoint** for this.

MemorySaver is the simplest checkpointer -- it stores graph state in Python memory (a dictionary). This is perfect for development and testing. In production, you would use a persistent backend like **SqliteSaver** or **PostgresSaver**.

Key ideas:

- Each graph execution is identified by a **thread_id** (a string you choose)
- The checkpointer saves the state at each step, so the graph can resume exactly where it left off
- Different **thread_id** values are completely independent -- like separate browser tabs

```
checkpointer = MemorySaver()
graph = workflow.compile(checkpointer=checkpointer)

# Each invocation needs a thread_id
config = {"configurable": {"thread_id": "investigation-001"}}
result = graph.invoke(initial_state, config=config)
```

Lab


Please complete Lab 13.4: Human in the Loop

Part 5: Multi-Agent Systems

Multi Agent Systems

- In the previous examples, we built agents which had rigid, pre-defined paths.
- The next evolution introduces when the AI can decide the path and the steps taken.
- To do this, we'll need a supervisor and more react agents.

Multi-Agent Systems



Multi Agent Systems

In previous labs, we built a single ReAct agent by hand -- defining nodes, routing functions, and tool execution logic ourselves. LangGraph provides `create_react_agent` as a **prebuilt shortcut** that packages all of that into a single function call.

`create_react_agent` takes:

- **model**: The LLM to use
- **tools**: A list of tools the agent can call (can be empty if the agent relies on LLM reasoning alone)
- **name**: A unique identifier used by the supervisor for delegation
- **prompt**: A system prompt that defines the agent's specialty and behavior

It returns a fully functional ReAct agent with:

- A tool-calling loop (reason -> act -> observe -> repeat)
- Automatic tool execution
- Message-based state management

Each specialist agent is a **self-contained expert** -- it knows its domain, has the right tools, and can work independently. The supervisor doesn't need to know *how* each agent works, only *what* each agent specializes in.

Multi Agent Systems

```
identity_analyst = create_react_agent(  
    model=llm,  
    tools=[check_device_reputation, check_login_velocity],  
    name="identity_analyst",  
    prompt="""You are an Identity Analyst specializing in account and session verification.
```

Your job is to:

1. Check the device reputation for this login using the device reputation tool
2. Review login velocity and MFA history using the login velocity tool
3. Assess whether the session looks like the legitimate account owner
4. Look for signs of credential stuffing or session hijacking

```
Provide a detailed analysis of your findings. Be specific about what you found and what it means for  
whether the account is compromised.""",  
)
```

The Supervisor

The `create_supervisor` function from `langgraph_supervisor` creates an agent that:

1. **Receives** the initial investigation request (an email to analyze)
2. **Decides** which specialist(s) to consult based on the email content
3. **Delegates** work by sending messages to specialist agents
4. **Receives** specialist reports back
5. **Synthesizes** all findings into a final verdict

Key parameters:

- **agents**: The list of specialist agents the supervisor can delegate to
- **prompt**: Instructions that tell the supervisor how to manage the team
- **output_mode**: Controls what the final output includes:
 - `"full_history"` -- returns the complete message trace (useful for debugging and observing handoffs)
 - `"last_message"` -- returns only the supervisor's final summary

The supervisor is itself an LLM-powered agent. It uses tool calls internally to "hand off" to specialists -- each specialist appears as a callable tool from the supervisor's perspective. This is the **multi-agent handoff** pattern.

The Supervisor

```
supervisor = create_supervisor(  
    model=llm,  
    agents=[identity_analyst, geo_network_analyst, behavior_analyst],  
    prompt="""You are the SOC Team Lead supervising a team of account-security analysts.
```

Your team:

- identity_analyst: Checks device reputation, login velocity, and MFA history
- geo_network_analyst: Checks geo-velocity / impossible travel and source IP reputation
- behavior_analyst: Analyzes the session's actions for anomalous behavior

For each login investigation:

1. Review the login event and determine which specialists to consult
2. Delegate to the relevant specialists (you may consult multiple)
3. After receiving their reports, synthesize the findings
4. Provide a final verdict: 'compromised', 'suspicious', or 'normal'
5. Include confidence level and key evidence from your team

Always consult at least 2 specialists for a thorough investigation. For logins from new locations, always consult the geo_network_analyst. For unusual devices, always consult the identity_analyst."""

```
    output_mode="full_history",  
)
```

```
# Compile the supervisor into a runnable graph  
graph = supervisor.compile()
```

Lab

Please complete Lab 13.5: Multi Agent Supervisor

Questions?